


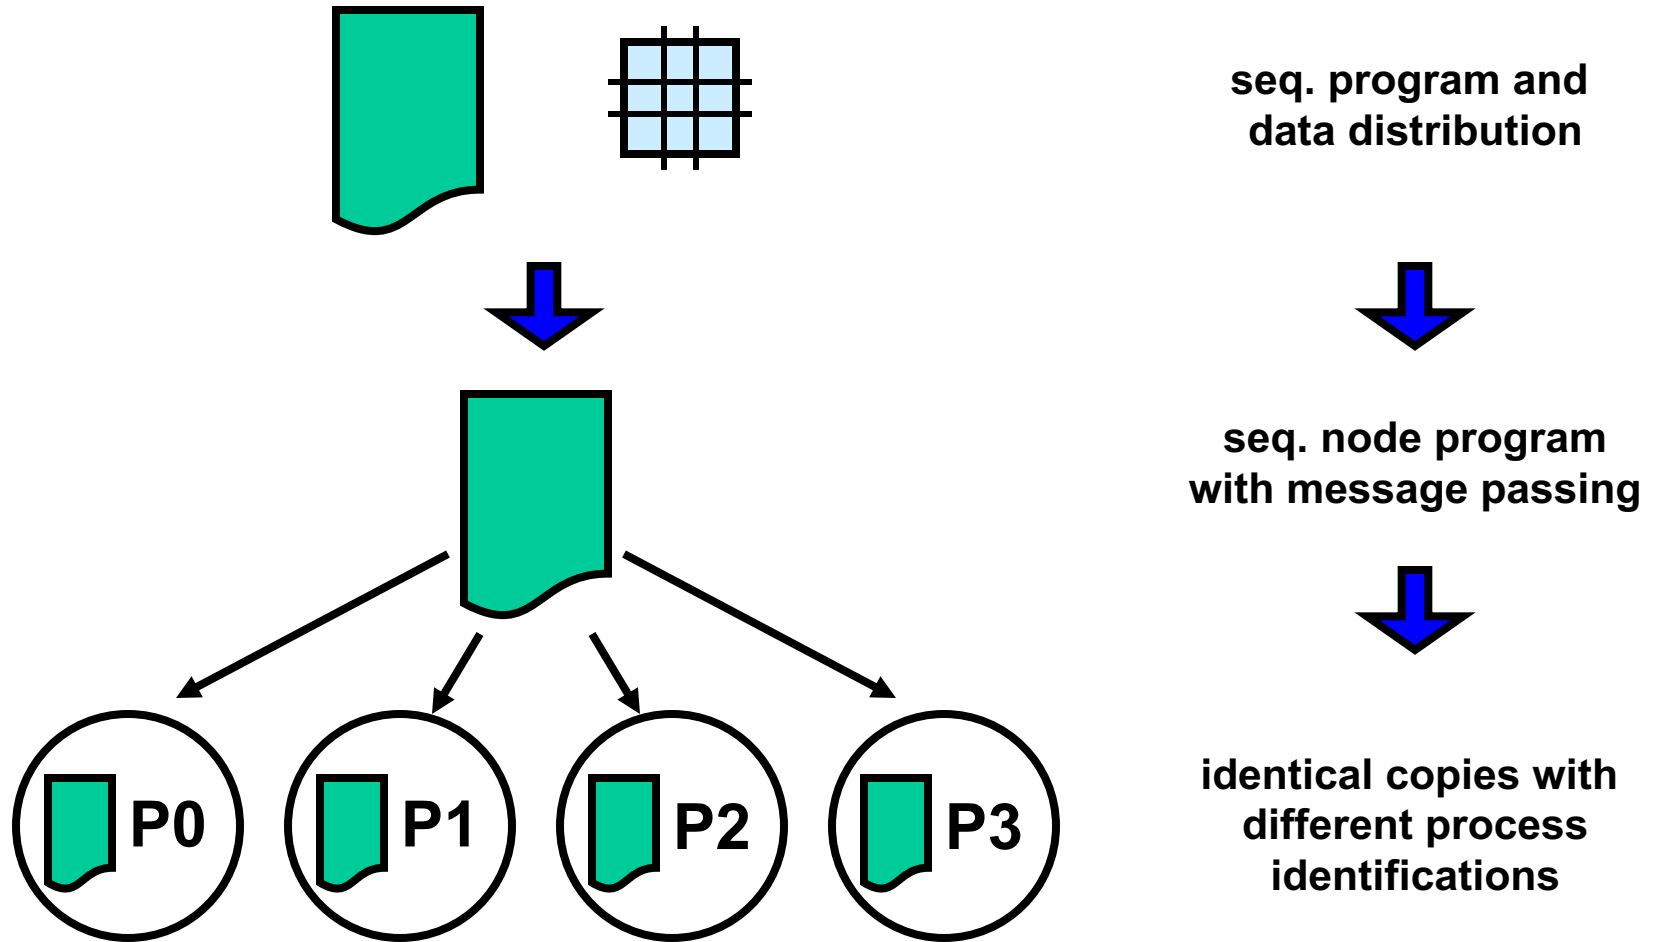
Introduction to the Parallelization with MPI

Michael Gerndt

Distributed Memory Architektur



Single Program Multiple Data (SPMD)



Program Parallelization

1. Adaptation of array declarations

- Local size of distributed arrays covers only the part of the data structure assigned to the process.

2. Index transformation

- Global indices are mapped into a tuple of node number and of a local index.

3. Work distribution

- Computations are executed by the process owning the assigned variable.

4. Communication

- Accesses to array elements of other processes have to be implemented by message passing.

Scope of the Message Passing Interface

- MPI 1.2
 - Point-to-Point communication
 - Collective communication
 - Communicators
 - Process topologies
 - User-defined data types
 - Operations and properties of the execution environment
 - Profiling interface
- MPI 2.0
 - Dynamic process creation
 - One-sided communication
 - Parallel IO
- MPI 3.0 (2012)
 - Nonblocking collectives
 - Additional one-sided communication operations
- MPI 4.0 Draft (2020)
- <https://www.mpi-forum.org>

Core Routines

- MPI 1.2 has 129 functions
- It is possible to write real programs with only six functions:
 - MPI_Init
 - MPI_Finalize
 - MPI_Comm_size
 - MPI_Comm_rank
 - MPI_Send
 - MPI_Recv

MPI_Init

```
int MPI_Init (int *argc, char ***argv)
```

IN argc, argv: *arguments*

return: *MPI_SUCCESS or error codes*

- This routine has to be called by each MPI process before any other MPI routine is executed
- Fortran interface
 - MPI_INIT (integer ierror)
 - The name is written in capital letters and the error code is returned via an additional argument.

MPI_Finalize

```
int MPI_Finalize ()
```

- Each process must call MPI_FINALIZE before it exits.
 - Precondition: All pending communication has to be finished.
 - One MPI_FINALIZE returns, no further MPI routines can be executed.
 - MPI_FINALIZE frees any resources.

MPI_Comm_size

```
int MPI_Comm_size (MPI_Comm comm, int *size)
```

IN comm: *Communicator*

OUT size: *Cardinality of the process group*

- **Communicator**

- Identifies a process group and defines the communication context. All message tags are unique with respect to a communicator.

- **MPI_COMM_WORLD**

- This is a predefined standard communicator. Its process group includes all processes of a parallel application.

- **MPI_Comm_size**

- It returns the number of processes in the process group of the given communicator.

MPI_Comm_rank

```
int MPI_Comm_rank (MPI_Comm comm, int *rank)
```

IN comm: *Communicator*

OUT rank: *process number of the executing process*

- **Process number**
 - The process number is a unique identifier within the process group of the communicator.
 - It is the only way to distinguish processes and to implement an SPMD program.
- **MPI_Comm_rank returns the process number of the executing process.**

MPI_Send

```
int MPI_Send (void *buf, int count, MPI_Datatype dtype, int dest, int tag,  
              MPI_Comm comm)
```

| | |
|-----------|-----------------------------------|
| IN buf: | <i>Address of the send buffer</i> |
| IN count: | <i>Number of data to be sent</i> |
| IN dtype: | <i>Data type</i> |
| IN dest: | <i>Receiver</i> |
| IN tag: | <i>Message tag</i> |
| IN comm: | <i>Communicator</i> |

- **MPI_Send**
 - Sends the data to the receiver.
 - It is a blocking operation, i.e. it terminates when the send buffer can be reused, either because the message was delivered or the data were copied to a system buffer.

MPI Data Types

C

| | |
|-------------------|------------------|
| MPI_CHAR | signed char |
| MPI_SHORT | signed short int |
| MPI_INT | signed int |
| MPI_LONG | signed long int |
| MPI_UNSIGNED_CHAR | |
| MPI_UNSIGNED_INT | |
| ... | |
| MPI_FLOAT | float |
| MPI_DOUBLE | double |
| MPI_LONG_DOUBLE | long double |
| MPI_BYTE | |
| MPI_PACKED | |

FORTRAN

| | |
|----------------------|------------------|
| MPI_INTEGER | integer |
| MPI_REAL | real |
| MPI_DOUBLE_PRECISION | double precision |
| MPI_COMPLEX | complex |
| MPI_LOGICAL | logical |
| MPI_CHARACTER | character(1) |
| MPI_BYTE | |
| MPI_PACKED | |

MPI_Recv

```
int MPI_Recv (void *buf, int count, MPI_Datatype dtype, int source, int  
              tag, MPI_Comm comm, MPI_Status *status)
```

OUT buf: *Address of the receive buffer*

IN count: *Size of receive buffer*

IN dtype: *Data type*

IN source: *Sender*

IN tag: *Message tag*

IN comm: *Communicator*

OUT status: *Status information*

- **Properties:**

- It is a blocking operation, i.e. it terminates after the message is available in the receive buffer.
- The message must not be larger than the receive buffer.
- The remaining part of the buffer not used for the received message will be unchanged.

Properties of MPI_Recv

- **Message selection**

- A message to be received by this function must match
 - the sender
 - the tag
 - the communicator
- Sender and tag can be specified as wild cards
 - MPI_ANY_SOURCE and MPI_ANY_TAG
- There is no wild card for the communicator.

- **Status**

- The data structure MPI_Status includes
 - status(MPI_SOURCE): sender of the message
 - status(MPI_TAG): message tag
 - status(MPI_ERROR): error code
- The actual length of the received message can be determined via MPI_Get_count.

Circular Left Shift Application

```
mpirun -np 4 shifts <number of positions>
```

Description

- Position 0 of an array with 100 entries is initialized to 1. The array is distributed among all processes in a blockwise fashion.
- A number of circular left shift operations is executed.
- The number is specified via a command line parameter.

[illegible]

Shifts: Initialization

```
#include "mpi.h"

main (int argc, char *argv[]) {
    int myid, np, ierr, lnbr, rnbr, shifts, i, j;
    int *values;
    MPI_Status status;

    ierr = MPI_Init (&argc, &argv);
    if (ierr != MPI_SUCCESS) {
        ...
    }

    MPI_Comm_size(MPI_COMM_WORLD, &np);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
```


Shifts: Definition of Neighbors


```
if (myid==0){
    lnbr=np-1; rnbr=myid+1;
}
else if (myid==np-1){
    lnbr=myid-1; rnbr=0;
}
else{
    lnbr=myid-1; rnbr=myid+1;
}

if (myid==0) shifts=atoi(argv[1]);
MPI_Bcast (&shifts, 1, MPI_INT, 0, MPI_COMM_WORLD);

values= (int *) calloc(100/np,sizeof(int));
if (myid==0){
    values[0]=1;
}
```

Shifts: Shift the array

```
for (i=0;i<shifts;i++){  
    int buf;  
  
    MPI_Send(&values[0],1,MPI_INT,lnbr,10,MPI_COMM_WORLD);  
    MPI_Recv(&buf, 1, MPI_INT,rnbr,10,  
            MPI_COMM_WORLD, &status);  
  
    for (j=1;j<100/np;j++){  
        values[j-1]=values[j];  
    }  
    values[100/np-1]=buf;  
}
```



MPI_Sendrecv

```
int MPI_Sendrecv (void *sendbuf, int sendcount, MPI_Datatype sendtype,  
                  int dest, int sendtag, void *recvbuf, int recvcount,  
                  MPI_Datatype recvtype, int source, MPI_Datatype recvtag,  
                  MPI_Comm comm, MPI_Status *status)
```

- Sendbuf and recvbuf have to be different
- Equivalent to the execution of MPI_Send and MPI_Receive in parallel threads.
- MPI_Sendrecv_replace:
 - Only one buffer.
 - The sent message is replaced by the received message.

Shifts: Shift the array

```
for (i=0;i<shifts;i++){
    int buf=values[0];

    for (j=1;j<100/np;j++){
        values[j-1]=values[j];
    }

    MPI_Sendrecv(&buf, 1, MPI_INT, lnbr, 10,
                 &values[100/np-1], 1, MPI_INT, rnbr, 10,
                 MPI_COMM_WORLD, &status);
}
```

MPI_Wtime

```
double MPI_Wtime (void)
```

- Return value represents elapsed wall-clock time since some time in the past measured in seconds.
- The time returned is local to the node that called it.

Collective Operations

- **Properties**

- Must be executed by all processes of the process group.
- Must be executed in the same sequence.
- All collective operations are blocking operations.

- **MPI provides three classes of collective operations**

- Synchronization
 - Barrier
- Communication
 - Broadcast
 - gather
 - scatter
- Reduction
 - Global value returned to one or all processes.
 - Combination with subsequent scatter.
 - Parallel prefix operations